200 XP

Knowledge check

2 minutes

1. You've trained a model using the Python SDK for Azure Machine Learning. You want to deploy the model as a containerized real-time service with high scalability and security. What kind of compute should you create to host the service?

- (●) An Azure Kubernetes Services (AKS) inferencing cluster ✓

  That is correct. You should use an AKS cluster to deploy a model as a scalable, secure, containerized service.

- ( ) A compute instance with GPUs
- ( ) A training cluster with multiple nodes.

2. You're deploying a model as a real-time inferencing service. What functions must the entry script for the service include?

- ( ) main() and score()
- ( ) base() and train()
- (●) init() and run() ✓

  That is correct. You must implement init and run functions in the entry (scoring) script.

Next unit: Summary

Continue