

Creating a batch inference pipeline

5 minutes

To create a batch inferencing pipeline, perform the following tasks:

1. Register a model

To use a trained model in a batch inferencing pipeline, you must register it in your Azure Machine Learning workspace.

To register a model from a local file, you can use the **register** method of the **Model** object as shown in the following example code:

Python

 Copy

```
from azureml.core import Model

classification_model = Model.register(workspace=your_workspace,
                                     model_name='classification_model',
                                     model_path='model.pkl', # local path
                                     description='A classification model')
```

Alternatively, if you have a reference to the **Run** used to train the model, you can use its **register_model** method as shown in the following example code:

Python

 Copy

```
run.register_model(model_name='classification_model',
                  model_path='outputs/model.pkl', # run outputs path
                  description='A classification model')
```

2. Create a scoring script

Batch inferencing service requires a scoring script to load the model and use it to predict new values. It must include two functions:

- **init()**: Called when the pipeline is initialized.
- **run(mini_batch)**: Called for each batch of data to be processed.

Typically, you use the `init` function to load the model from the model registry, and use the `run` function to generate predictions from each batch of data and return the results. The following example script shows this pattern:

Python

 Copy

```
import os
import numpy as np
from azureml.core import Model
import joblib

def init():
    # Runs when the pipeline step is initialized
    global model

    # load the model
    model_path = Model.get_model_path('classification_model')
    model = joblib.load(model_path)

def run(mini_batch):
    # This runs for each batch
    resultList = []

    # process each file in the batch
    for f in mini_batch:
        # Read comma-delimited data into an array
        data = np.genfromtxt(f, delimiter=',')
        # Reshape into a 2-dimensional array for model input
        prediction = model.predict(data.reshape(1, -1))
        # Append prediction to results
        resultList.append("{}: {}".format(os.path.basename(f), prediction[0]))
    return resultList
```

3. Create a pipeline with a `ParallelRunStep`

Azure Machine Learning provides a type of pipeline step specifically for performing parallel batch inferencing. Using the `ParallelRunStep` class, you can read batches of files from a **File** dataset and write the processing output to a **OutputFileDatasetConfig**. Additionally, you can set the **output_action** setting for the step to "append_row", which will ensure that all instances of the step being run in parallel will collate their results to a single output file named *parallel_run_step.txt*:

Python

 Copy

```
from azureml.pipeline.steps import ParallelRunConfig, ParallelRunStep
from azureml.data import OutputFileDatasetConfig
from azureml.pipeline.core import Pipeline
```

```
# Get the batch dataset for input
batch_data_set = ws.datasets['batch-data']

# Set the output location
output_dir = OutputFileDatasetConfig(name='inferences')


# Define the parallel run step configuration
parallel_run_config = ParallelRunConfig(
    source_directory='batch_scripts',
    entry_script="batch_scoring_script.py",
    mini_batch_size="5",
    error_threshold=10,
    output_action="append_row",
    environment=batch_env,
    compute_target=aml_cluster,
    node_count=4)

# Create the parallel run step
parallelrun_step = ParallelRunStep(
    name='batch-score',
    parallel_run_config=parallel_run_config,
    inputs=[batch_data_set.as_named_input('batch_data')],
    output=output_dir,
    arguments=[],
    allow_reuse=True
)

# Create the pipeline
pipeline = Pipeline(workspace=ws, steps=[parallelrun_step])
```

4. Run the pipeline and retrieve the step output

After your pipeline has been defined, you can run it and wait for it to complete. Then you can retrieve the **parallel_run_step.txt** file from the output of the step to view the results, as shown in the following code example:

Python	 Copy
<pre>from azureml.core import Experiment # Run the pipeline as an experiment pipeline_run = Experiment(ws, 'batch_prediction_pipeline').submit(pipeline) pipeline_run.wait_for_completion(show_output=True) # Get the outputs from the first (and only) step prediction_run = next(pipeline_run.get_children()) prediction_output = prediction_run.get_output_data('inferences') prediction_output.download(local_path='results') # Find the parallel_run_step.txt file for root, dirs, files in os.walk('results'): for file in files:</pre>	

```
if file.endswith('parallel_run_step.txt'):
    result_file = os.path.join(root,file)

# Load and display the results
df = pd.read_csv(result_file, delimiter=":", header=None)
df.columns = ["File", "Prediction"]
print(df)
```

Next unit: Publishing a batch inference pipeline

Continue >