

Introduction

Azure Notebooks is a cloud-hosted service that enables users to develop and run Jupyter notebooks with no installation. Jupyter is an open-source project that combines Markdown text, graphical visualizations, and code into a single platform. It has become a popular tool for many use cases such as data science, data cleaning, and statistical modeling.


In this lab step, you will create an Azure Notebook in Azure Machine Learning Studio from an existing notebook on GitHub.

Instructions


1. Right-click the following [link](#) and select **Save link as**:

Save to a location on your computer, this Jupyter notebook file will be used in a later step.

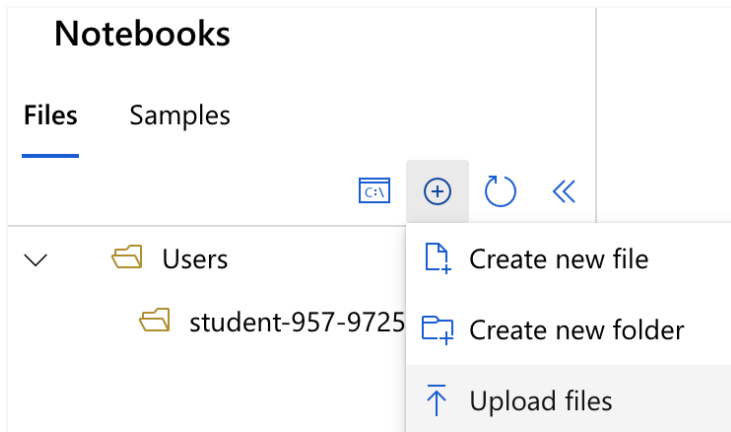
2. In the Azure Machine Learning Studio, click on **Notebooks** on the left-hand side:



3. Under **Quick Actions**, click **Upload an existing notebook or file**:



If you don't see the notebooks welcome screen, you can instead click the **+** icon followed by **Upload files**:

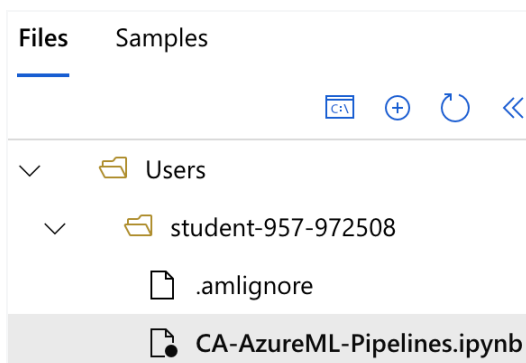


4. In the **Upload files** dialog that appears, perform the following before clicking **Upload**:

- Click **Click to browse and select file(s)** and locate the Jupyter notebook file saved earlier in this lab step
- Check the **I trust contents of this file**


The Jupyter notebook will be uploaded to the Notebook user's folder.

5. Click the **CA-AzureML-Pipelines.ipynb** file to open up the notebook:



You will see the Jupyter notebook load to the right, you will work through this notebook later.

6. In the top bar, under the **Compute** section, Click the **+** button to create a new compute instance:



The Azure Notebook will run the Jupyter file within a compute instance. So one must be created in order to use the notebook file.

7. Configure the **required settings** for the compute instance, leaving the rest of the settings as defaults before clicking **Create**:

- **Compute name:** *ca-#####-compute*, replacing the ##### characters with unique numbers.

Create compute instance

Required Settings

Advanced Settings

Configure required settings

Select the name and virtual machine size you would like to use for your compute instance. Please note that a compute instance can not be shared. It can only be used by a single assigned user. By default, it will be assigned to the creator and you can change this to a different user in the advanced settings section.

Compute name *

ca-41412-compute

Location

westcentralus

Virtual machine type

CPU GPU

Virtual machine size

Select from recommended options Select from all options

Total available quota: 50 cores

Name ↑	Category	Workload types	Av...	Cost
Standard_DS11_v2 2 cores, 14GB RAM, 28GB storage	Memory optimized	Development on Notebooks (or other IDE) and light weight testing	50 co...	\$0.17/hr
Standard_DS3_v2 4 cores, 14GB RAM, 28GB storage	General purpose	Classical ML model training, AutoML runs, pipeline runs (default compute)	50 co...	\$0.25/hr
Standard_DS12_v2 4 cores, 28GB RAM, 56GB storage	Memory optimized	Training on large datasets (>1GB) parallel run steps, batch inferencing	50 co...	\$0.33/hr
Standard_F4s_v2 4 cores, 8GB RAM, 32GB storage	Compute optimized	Real-time inferencing and other latency-sensitive tasks	10 co...	\$0.20/hr

The compute instance name needs to be globally unique, which is why you entered the unique numbers in the name.

Note: The compute instance will start provisioning and will take about 5 minutes. The status will be shown in the compute dropdown.

Compute:

ca-41412-compute - Creating

Once the status is set to running, the notebook can be used:

<https://cloudacademy.com/lab/building-azure-ml-pipelines-using-azure-machine-learning-sdk/creating-an-azure-notebook-for-the-azure-ml-pipeline...> 3/4



Summary

In this lab step, you created an Azure Notebook by importing a Jupyter notebook and creating a compute instance for hosting the notebook.